


Introduction to Dynamic Programming: Basic theory and numerical tools

QuantEcon-RSE Honours workshop 2018

Fedor Iskhakov, Australian National University


QuantEcon-RSE
March 10, 2018

Value of an infinite stream of payments



- Interest rate $r > 0$
- What is the value of the annuity V ?

Value of an infinite stream of payments



- Interest rate $r > 0$
- What is the value of the annuity V ?
- Discounted present values

$$\frac{c}{(1+r)^0} \quad \frac{c}{(1+r)^1} \quad \frac{c}{(1+r)^2} \quad \dots$$

- Let $\beta = \frac{1}{1+r}$

$$V = c + \beta c + \beta^2 c + \dots = \sum_{t=0}^{\infty} \beta^t c$$

Value of an infinite stream of payments

- Note that $\beta = \frac{1}{1+r} < 1$ because $r > 0 \Rightarrow$ converging geometric series

$$V = c + \beta c + \beta^2 c + \dots = \sum_{t=0}^{\infty} \beta^t c = \frac{c}{1 - \beta}$$

- Another approach: reformulate this as a **recursive equation**

$$\begin{aligned} V &= c + \beta(c + \beta^2 c + \dots) \\ &= c + \beta V \end{aligned}$$

- Extremely simple case here, leads to the same answer
- In general, we need a range of computational tools which we can first try on this simple problem

Backward induction

- 1 Start with a guess V_0
- 2 Insert into the recursive equation $V_1 = c + \beta V_0$
- 3 Insert new value V_1 into the recursive equation again $V_2 = c + \beta V_1$
- 4 Repeat until **convergence** (i denotes iteration number)

$$\|V_i - V_{i-1}\| \leq \varepsilon \text{ (small number)}$$

- Will this converge?

$$\|V_i - V_{i-1}\| = \|(c + \beta V_{i-1}) - (c + \beta V_{i-2})\| = \beta \|V_{i-1} - V_{i-2}\|$$

- $\beta < 1 \Rightarrow$ with every iteration i the difference $\|V_i - V_{i-1}\|$ becomes smaller (recursive formula is a **contraction mapping**)
- Banach fixed point theorem guarantees unique solution!

Backward induction

- 1 Start with a guess V_0
- 2 Insert into the recursive equation $V_1 = c + \beta V_0$
- 3 Insert new value V_1 into the recursive equation again $V_2 = c + \beta V_1$
- 4 Repeat until **convergence** (i denotes iteration number)

$$||V_i - V_{i-1}|| \leq \varepsilon \text{ (small number)}$$

- Will this converge?

$$||V_i - V_{i-1}|| = ||(c + \beta V_{i-1}) - (c + \beta V_{i-2})|| = \beta ||V_{i-1} - V_{i-2}||$$

- $\beta < 1 \Rightarrow$ with every iteration i the difference $||V_i - V_{i-1}||$ becomes smaller (recursive formula is a **contraction mapping**)
- Banach fixed point theorem guarantees unique solution!

Backward induction

- 1 Start with a guess V_0
- 2 Insert into the recursive equation $V_1 = c + \beta V_0$
- 3 Insert new value V_1 into the recursive equation again $V_2 = c + \beta V_1$
- 4 Repeat until **convergence** (i denotes iteration number)

$$||V_i - V_{i-1}|| \leq \varepsilon \text{ (small number)}$$

- Will this converge?

$$||V_i - V_{i-1}|| = ||(c + \beta V_{i-1}) - (c + \beta V_{i-2})|| = \beta ||V_{i-1} - V_{i-2}||$$

- $\beta < 1 \Rightarrow$ with every iteration i the difference $||V_i - V_{i-1}||$ becomes smaller (recursive formula is a **contraction mapping**)
- Banach fixed point theorem guarantees unique solution!

Backward induction

- 1 Start with a guess V_0
- 2 Insert into the recursive equation $V_1 = c + \beta V_0$
- 3 Insert new value V_1 into the recursive equation again $V_2 = c + \beta V_1$
- 4 Repeat until **convergence** (i denotes iteration number)

$$||V_i - V_{i-1}|| \leq \varepsilon \text{ (small number)}$$

- Will this converge?

$$||V_i - V_{i-1}|| = ||(c + \beta V_{i-1}) - (c + \beta V_{i-2})|| = \beta ||V_{i-1} - V_{i-2}||$$

- $\beta < 1 \Rightarrow$ with every iteration i the difference $||V_i - V_{i-1}||$ becomes smaller (recursive formula is a **contraction mapping**)
- Banach fixed point theorem guarantees unique solution!

Backward induction

- 1 Start with a guess V_0
- 2 Insert into the recursive equation $V_1 = c + \beta V_0$
- 3 Insert new value V_1 into the recursive equation again $V_2 = c + \beta V_1$
- 4 Repeat until **convergence** (i denotes iteration number)

$$\|V_i - V_{i-1}\| \leq \varepsilon \text{ (small number)}$$

- Will this converge?

$$\|V_i - V_{i-1}\| = \|(c + \beta V_{i-1}) - (c + \beta V_{i-2})\| = \beta \|V_{i-1} - V_{i-2}\|$$

- $\beta < 1 \Rightarrow$ with every iteration i the difference $\|V_i - V_{i-1}\|$ becomes smaller (recursive formula is a **contraction mapping**)
- Banach fixed point theorem guarantees unique solution!

Numerical illustration

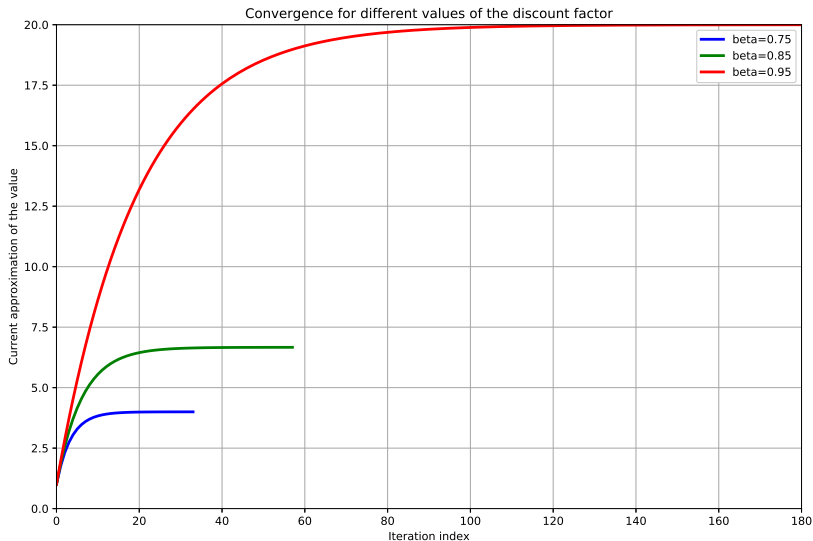
The code:
annuity.ipynb



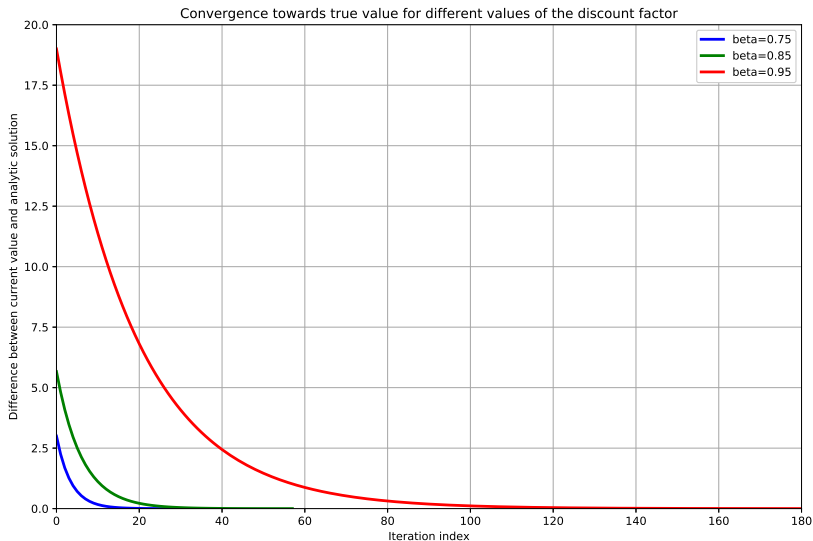
Questions and discussion

- 1 Convergence to the analytical solution?
- 2 What determines the rate of convergence?
- 3 What happens when $\beta = 1$?

Convergence speed and the role of β



Convergence speed and the role of β



Cake eating problem



- Cake of initial size W_0
- How much of the cake to eat each period t ?
- What is not eaten in period t is left for the future

$$W_{t+1} = W_t - c_t$$

- Utility flow from cake consumption

$$u(c_t) = \log(c_t)$$

- Future is discounted with discount factor β
- Optimization problem

$$\max_{\{c_t\}_0^\infty} \sum_{t=0}^{\infty} \beta^t u(c_t) \longrightarrow \max$$

Recursive formulation = Bellman equation

- **Value function** $V(W_t)$ = the maximum attainable value given the size of cake W_t (in period t)

$$\begin{aligned} V(W_0) &= \max_{\{c_t\}_0^\infty} \sum_{t=0}^{\infty} \beta^t u(c_t) \\ &= \max_{c_0} \{u(c_0) + \beta \max_{\{c_t\}_1^\infty} \sum_{t=1}^{\infty} \beta^{t-1} u(c_t)\} \\ &= \max_{c_0} \{u(c_0) + \beta V(W_1)\} \end{aligned}$$

- The Bellman equation is

$$V(W_t) = \max_{0 \leq c_t \leq W_t} \left\{ u(c_t) + \beta \underbrace{V(W_{t+1})}_{=W_t - c_t} \right\}$$

Dynamic programming

“DP is a recursive method for solving sequential decision problems”



John Rust 2006 *New Palgrave Dictionary of Economics*

State variables vector of variables that describe all relevant information about the modeled decision process, W_t

Decision variables vector of variables describing the choices, c_t

Instantaneous payoff utility function, $u(c_t)$, with time separable discounted utility

Timing scale discrete, finite/infinite horizon, discount factor β

Motion rules agent's beliefs of how state variable evolve through time, conditional on choices, $W_{t+1} = W_t - c_t$

Value function maximum attainable utility, function of the state variables, $V(W_t)$

Policy function mapping from state space to action space that returns the optimal choice, $c^*(W_t)$

Cake Eating: Analytical Solution

Guess and verify

- Start with a (good) guess of $V(W) = A + B \log W$

$$\begin{aligned} V(W) &= \max_c \{u(c) + \beta V(W - c)\} \\ A + B \log W &= \max_c \{ \log c + \beta(A + B \log(W - c)) \} \end{aligned}$$

Exercise: Determine A and B and find the optimal rule for cake consumption.

- This is only possible in *few* models!
- For cake eating problem

$$c^*(W) = \arg \max_c \{ \log(c) + \beta V(W - c) \} = (1 - \beta)W$$

Cake Eating: Numerical Solution

Will **backward induction** work as before?

- Value of annuity:

$$V = c + \beta V$$

- Cake eating:

$$V(W) = \max_{0 \leq c \leq W} \{u(c) + \beta V(W - c)\}$$

- Have to solve the **functional equation** for $V(W)$
- The Bellman operator in functional space

$$T(V)(W) \equiv \max_{0 \leq c \leq W} \{u(c) + \beta V(W - c)\}$$

- The Bellman equations is then $V(W) = T(V)(W)$, with the solution given by the fixed point

Can we find the fixed point by iterations?

- Need contraction property for $T(V)(W)$
- Blackwell sufficient conditions for contraction
 - ① Monotonicity: satisfied due to maximization in $T(V)(W)$
 - ② Discounting: satisfied by elementary argument when $\beta < 1$
- The Bellman operator is a **contraction mapping**!

Contraction Mapping Theorem (Banach Fixed Point Theorem)

Let (S, ρ) be a complete metric space with a contraction mapping $T : S \rightarrow S$. Then

- ① T admits a unique fixed-point $V^* \in S$, i.e. $T(V^*) = V^*$.
- ② V^* can be found by repeated application of the operator T , i.e. $T^n(V) \rightarrow V^*$ as $n \rightarrow \infty$.

Value function iterations (VFI)

- 1 Start with an arbitrary guess $V_0(W)$
- 2 At each iteration i compute

$$\begin{aligned} V_i(W) = T(V_{i-1})(W) &= \max_{0 \leq c \leq W} \{u(c) + \beta V_{i-1}(W - c)\} \\ c_{i-1}(W) &= \arg \max_{0 \leq c \leq W} \{u(c) + \beta V_{i-1}(W - c)\} \end{aligned}$$

- 3 Repeat until convergence

$$\|V_i(W) - V_{i-1}(W)\| \leq \varepsilon \text{ (small number, } \|\cdot\| \text{ sup norm)}$$

The contraction mapping theorem implies:

- Unique fixed point \Leftrightarrow unique solution to the Bellman equation
- The fixed point can be reached by an iterative process using an **arbitrary initial guess!**
- Therefore VFI algorithm converges globally

Cake Eating: numerical implementation

How to numerical implement the Bellman operator?

- Cake is continuous and thus value function is a function of continuous variable

- Solution: **discretize** W

Construct a *grid* (vector) of cake-sizes $\vec{W} \in \{0, \dots, \bar{W}\}$

$$V_i(\vec{W}) = \max_{0 \leq c \leq \vec{W}} \{u(c) + \beta V_{i-1}(\vec{W} - c)\}$$

- Compute value and policy function sequentially point-by-point
- May need to compute the value function **between grid points**
 \Rightarrow Interpolation and function approximation

Cake Eating: decision-state grid

Can interpolation be avoided?

- Note that conditional on W_t , the choice of c defines W_{t+1}
- Can replace c with W_{t+1} in Bellman equation so that **next period cake size is the decision variable**
- “Dual” formulation of the same problem

$$V_i(\vec{W}) = \max_{0 \leq \vec{W}' \leq \vec{W}} \{u(\vec{W} - \vec{W}') + \beta V_{i-1}(\vec{W}')\}$$

- Compute value and policy function sequentially point-by-point
- Note that grid $\vec{W} \in \{0, \dots, \bar{W}\}$ is used twice: for state space and for decision space
- Only precise when number of grid points is large

Cake eating: Numerical implementation I

The code:
cake1.ipynb



Questions and discussion

- 1 Rate of convergence?
- 2 Magnitude of numerical errors?
- 3 The role of grid density?

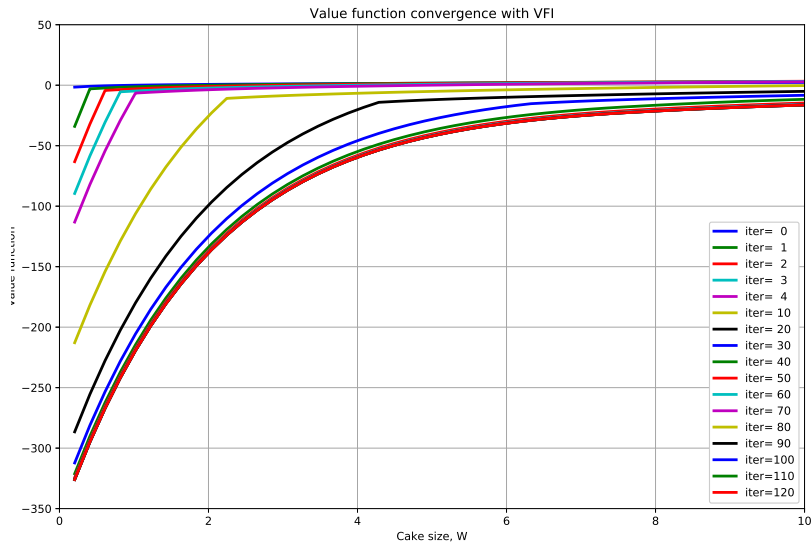
How to measure numerical errors?

- In our case there is an analytic solution

$$c^*(W) = (1 - \beta)W$$

- Typically very dense (slow) grid is used in place of true solution
- Can control for max or mean error at the grid points of value and policy functions

Computed value functions



Cake Eating: another numerical implementation

Control for grid over state space separately from the discretization of the choice variables to increase accuracy

- As before solve cake eating Bellman equation by VFI

$$V(W) = \max_{0 \leq c \leq W} \{u(c) + \beta V(W - c)\}$$


- Discretize state space with $\vec{W} \in \{0, \dots, \bar{W}\}$
- Discretize decision space with $\vec{D} \in \{0, \dots, \bar{D}\}$, usually $\bar{D} = \bar{W}$

$$V_i(\vec{W}) = \max_{0 \leq \vec{D} \leq \vec{W}} \{u(c) + \beta V_{i-1}(\vec{W} - c)\}$$

- Compute value/policy function point-by-point on grid \vec{W}
- Find the maximum over the points of grid \vec{D} that satisfy the choice set condition $0 \leq \vec{D} \leq W$
- Now have to compute the value function **between grid points**

Function interpolation

- Cake consumption can take on infinitely many cake sizes
→ we need to be able to approximate the value function for all resulting levels of cake left to subsequent periods, $W' = W - c$.
- We have a grid \vec{W} , a set of points for which we have explicitly found the value function, $V(\vec{W})$



- For now, assume interpolation function is $\check{V}(W)$

Cake eating: Numerical implementation II


The code:
cake2.ipynb



Questions and discussion

- 1 Did convergence speed change?
- 2 Did magnitude of numerical errors change?

Computed value functions



Cake Eating: continuous choice implementation

Is it possible not to discretize the choice variable? Yes!

- Discretize state space with $\vec{W} \in \{0, \dots, \bar{W}\}$

$$V_i(\vec{W}) = \max_{0 \leq c \leq \vec{W}} \{u(c) + \beta V_{i-1}(\vec{W} - c)\}$$

- Compute value/policy function point-by-point on grid \vec{W}
- Apply continuous solver to find the solution of the maximization problem inside Bellman equation for each point of the grid \vec{W}

OR

- Solve the first order conditions given by the **Eurler equation**
- Contraction mapping property does not rely in the nature of the control variable \Rightarrow VFI still works!

Euler Equation

- ① First order condition (FOC) of the Bellman equation w.r.t. c

$$0 = u'(c_t) + \beta \frac{\partial V(W_{t+1})}{\partial c_t} = u'(c_t) + \beta V'(W_{t+1}) \underbrace{\frac{\partial W_{t+1}}{\partial c_t}}_{-1},$$

$$\Rightarrow u'(c_t^*) = \beta V'(W_{t+1})$$

- ② From the Envelope theorem we have

$$V'(W_t) = \frac{\partial}{\partial W_t} [u(c_t) + \beta V(W_{t+1})] \Big|_{c_t^*} = \beta V'(W_{t+1}) \underbrace{\frac{\partial W_{t+1}}{\partial W_t}}_1$$

- ③ Combine and plug back into FOC

$$u'(c_t^*) = V'(W_t) \Rightarrow u'(c_{t+1}^*) = V'(W_{t+1}) \Rightarrow$$

$$u'(c_t^*) = \beta u'(c_{t+1}^*)$$

VFI with non-linear optimizer or solver

- 1 Discretize state space with $\vec{W} \in \{0, \dots, \bar{W}\}$
- 2 Start with an arbitrary guess $V_0(\vec{W})$ or $c_0(\vec{W})$
- 3 At each iteration i and for each point W_j on grid \vec{W} compute

$$V_i(W_j) = \max_{0 \leq c \leq W_j} \{u(c) + \beta V_{i-1}(W_j - c)\}$$

or solve

$$u'(c_i(W_j)) = \beta u'(c_{i-1}(W_j - c))$$

- 4 Repeat until convergence

$$\max_j (V_i(W_j) - V_{i-1}(W_j)) \leq \varepsilon \text{ (small number)}$$

$$\max_j (c_i(W) - c_{i-1}(W)) \leq \varepsilon$$

Cake eating: Numerical implementation III

The code:
Homework



Questions and discussion

- 1 Did convergence speed change?
- 2 Did magnitude of numerical errors change?
- 3 Any ideas about how to speed up the computations?

Only a primer in dynamic programming

Many topics and extensions left out:

1 Finite horizon formulation

- Time scripts essential $V_t(W_t)$ and $c_t(W_t)$
- Final period T with $c_T(W_T) = W_T$ and $V_T(W_T) = u(W_T)$
- Convergence replaced with backward calculation until $t = 0$

2 Stochastic models

- Stochastic evolution of state space, controlled by decisions:
random returns, controlled and exogenous Markov processes
- Idiosyncratic random components in preferences (taste shocks)
- Need for numerical integration in Bellman equation

Only a primer in dynamic programming

Many topics and extensions left out:

- ③ Higher dimension
 - Multiple state variables
 - Multi-dimensional decisions
 - Need for high power computing methods

- ④ Alternative (faster) solution approaches
 - Policy function iterations, faster convergence
 - Polyalgorithm using VFI and Newton-Raphson iteration for robust speed up
 - Endogenous gridpoint methods to avoid root-finding operations in models of applicable class

Thank you!

Questions?